

TVM BENCHMARK

2025-10-21 김규진

TVM BenchMark

```
class PyTorchModel(torch.nn.Module):
    def __init__(self):
        super(PyTorchModel, self).__init__()
        self.fc1 = torch.nn.Linear(784, 256)
        self.relu1 = torch.nn.ReLU()
        self.fc2 = torch.nn.Linear(256, 10, bias=False)

    def forward(self, x):
        x = self.fc1(x)
        x = self.relu1(x)
        x = self.fc2(x)
        return x
```

Tensor

```
=====
BENCHMARKING PYTORCH MODEL
=====
✓ Warming up PyTorch inference...
✓ Benchmarking PyTorch inference...

PyTorch Results (over 100 runs):
  Average: 0.02 ms
  Std Dev: 0.01 ms
  Min: 0.01 ms
  Max: 0.10 ms

=====
BENCHMARKING TVM MODEL
=====
✓ Warming up TVM inference...
✓ Benchmarking TVM inference...

TVM Results (over 100 runs):
  Average: 0.28 ms
  Std Dev: 0.05 ms
  Min: 0.25 ms
  Max: 0.63 ms
```

Resnet 18 - TVM

```
=====
TVM CPU-OPTIMIZED INFERENCE TIME (over 100 runs)
=====
Average: 1959.64 ms
Std Dev: 22.03 ms
Min: 1914.91 ms
Max: 2058.70 ms

=====
✓ Inference completed - Output shape: (1, 1000)

=====
TOP-5 PREDICTIONS FOR DOG IMAGE
=====
1. Pomeranian 39.73% 
2. Samoyed 22.64%
3. keeshond 6.05%
4. Japanese spaniel 5.42%
5. Arctic fox 4.58%
```

Resnet 18 - Torch

```
=====
PYTORCH INFERENCE TIME (over 100 runs)
=====
Average: 13.89 ms
Std Dev: 1.69 ms
Min: 11.18 ms
Max: 19.58 ms

=====
✓ Inference completed - Output shape: (1, 1000)

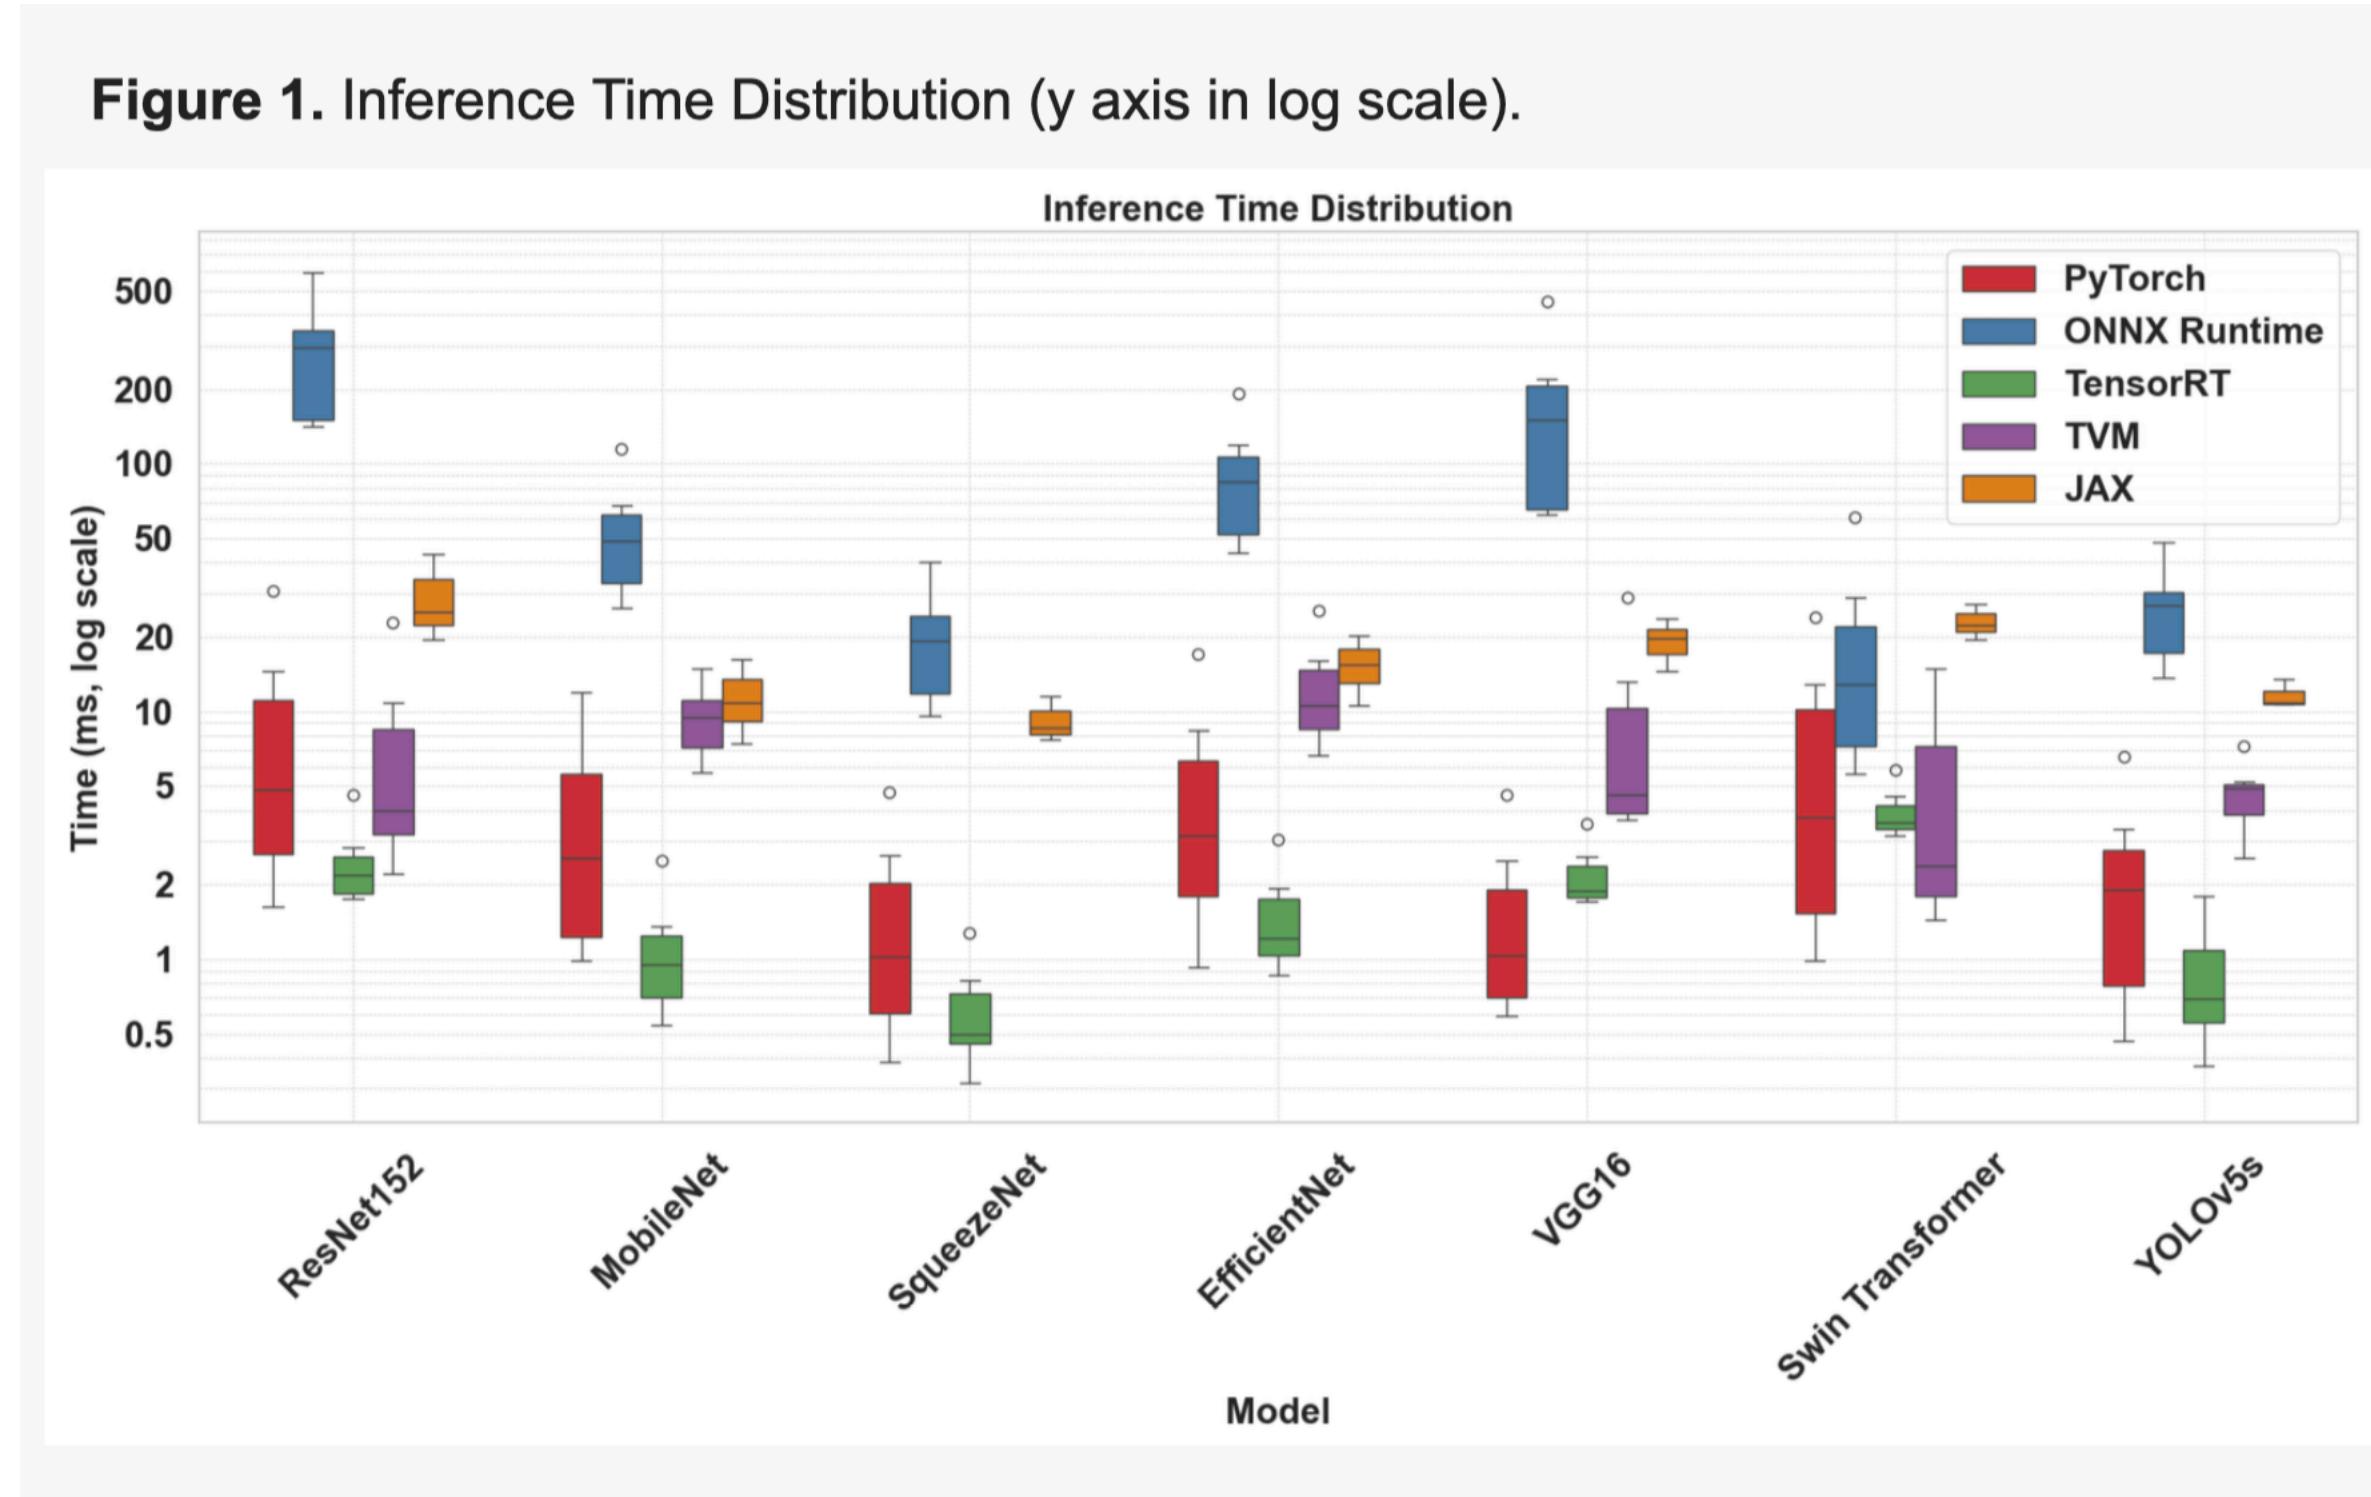
=====
TOP-5 PREDICTIONS FOR DOG IMAGE
=====
1. Pomeranian 39.73% 
2. Samoyed 22.64%
3. keeshond 6.05%
4. Japanese spaniel 5.42%
5. Arctic fox 4.58%
```

진행사항

- PyTorch: Input → C++ Kernel → Output
- TVM: Input → VM Interpreter → Compiled Function → Output
- Function OverHead? -> 딜레이에 미치는 요소가 큰 비율은 아님
- MetaSchedule (AutoTuning)? -> 이게 핵심인데 tvm relax 의 스케줄링이 모델 로그 스케일이 커짐에 따라 런타임 딜레이가 줄어들도록 설계하고있음 but 아직 0.22dev 버전이기 때문에 지원하지 않는 op가 많아 큰 모델을 테스트해볼 수가 없음
- Tensor IR 을 그대로 사용했을때 llvm core 스케줄링이 들어가도 torch 가 더 빠른게 의문

참고

- https://www.mdpi.com/2079-9292/14/15/2977?utm_source=chatgpt.com



다음 단계

- Tensor RT Pipeline
- TVM Virtual Machine Scheduling